

HW #4

Virtual Memory Simulator

Yunmin Go

School of CSEE



Requirements

- Implement a virtual memory simulator for multi-process
 - In this assignment, you will implement a virtual memory simulator.
 - Although the virtual memory simulator you will be implementing in this assignment is highly simplified, through this implementation, you will gain a general understanding of how an operating system executes processes.
 - Please refer to the skeleton code provided through the LMS along with the assignment requirements for implementation.
 - Please read the following requirements carefully.

Requirements

■ Program usage

Usage: `./vmsim <process image files... up to 5 files>`

- The simulator's name is 'vmsim'.
- The simulator can load highly simplified process images from text files.
- The maximum number of processes that the simulator can load is defined as 'MAX_PROCESSES' in 'vmsim.h'.
 - The default value for MAX_PROCESSES is 5.

Requirements

■ Process image

■ Example of 'proc0.txt'

Size of process image (This process can have 16384 bytes of memory space)

The diagram illustrates the structure of a process image file. It shows a list of memory operations (M, A, S, L) with their respective parameters. The first two numbers in each line represent the memory address and the number of bytes, respectively. The third number represents the instruction number. The first line is circled in red, and an arrow points to the text 'The number of instructions to execute (Simulator executes the following 9 instructions)'. Another arrow points to the first line of the list, indicating the start of the instructions to execute.

16384	9	
M	0	10
M	1	20
M	2	30
A	3	0 1
A	4	1 2
S	3	0x1000
S	4	0x2000
L	5	0x1000
L	6	0x2000

→ The number of instructions to execute
(Simulator executes the following 9 instructions)

→ Instructions to execute
In the simulator, instruction is stored as string in the pseudo physical memory.
Please refer to page 6 and 10.

- When the simulator starts, the process image (including string of instructions) provided by files are loaded into pseudo physical memory.

Requirements

■ Instructions

- The simulator can execute four types of instructions.

- 'M' type: Move integer value into register

- ex) M 0 10 → Move integer 10 to Register#0.
 ↑ ↑
 Reg Integer constant

- 'A' type: Add the integer value of RegSrc1 and the integer value of RegSrc2. The result is stored in RegDst.

- ex) A 3 0 1 → Register#3 stores the addition result 'value in Register#0 + value in Register#1'.
 ↑ ↑ ↑
 RegDst RegSrc1 RegSrc2

- 'L' type: Load 4-bytes integer value from the specified memory address into register.

- ex) L 5 0x1000 → Register#5 stores the 4-bytes integer value loaded from the (virtual) memory address 0x1000
 ↑ ↑
 Reg Memory address

- 'S' type: Store 4-bytes integer value of register at the specified memory address.

- ex) S 3 0x1000 → The 4-bytes integer value of Register#3 is stored at the (virtual) memory address 0x1000
 ↑ ↑
 Reg Memory address

Requirements

■ Instructions

- The simulator provides 8 pseudo registers which is declared as `'int register_set[MAX_REGISTERS]'` in `'vmsim.h'`.
 - The default value for `MAX_REGISTERS` is 8.
 - Each register has 4-bytes since it is declared as integer.
- In the simulator, each instruction is stored as string (char array) in the pseudo physical memory and its maximum string length is defined by `'INSTRUCTION_SIZE'` in `'vmsim.h'`.
 - The default value for `INSTRUCTION_SIZE` is 32 bytes.

Requirements

■ Instructions

- The simulator uses a program counter maintained by each process to fetch instruction strings from memory.
 - Please refer to 'int pc' in structure 'Process'.
 - In the simulator, the first instruction of a process is always located at 0x00000000 (virtual memory), so 'pc' is initially set to 0x00000000.
 - PC must increment by INSTRUCTION_SIZE to point to the next instruction.
- Instruction operations are already implemented in 'vmsim_op.c', so there is no need to modify 'vmsim_op.c'
 - `void op_move(Process *process, char *instruction);`
 - `void op_add(Process *process, char *instruction);`
 - `void op_load(Process *process, char *instruction);`
 - `void op_store(Process *process, char *instruction);`

Requirements

- Context switching
 - The simulator works with a pseudo clock.
 - We assume that the simulator executes any instruction within a single clock cycle.
 - The simulator performs context switching at every clock to simulate multi-process environment.
 - Process scheduling follows a round-robin scheme across all clocks.
 - Upon a context switch, the resumed process restores the saved register set from the temporary register set. Additionally, a process that completes an instruction stores its current register set in a temporary register set.
 - Please refer to 'int execute(Process *process)' in 'vmsim_main.c' and 'int temp_reg_set[MAX_REGISTERS]' in structure 'Process'.

Requirements

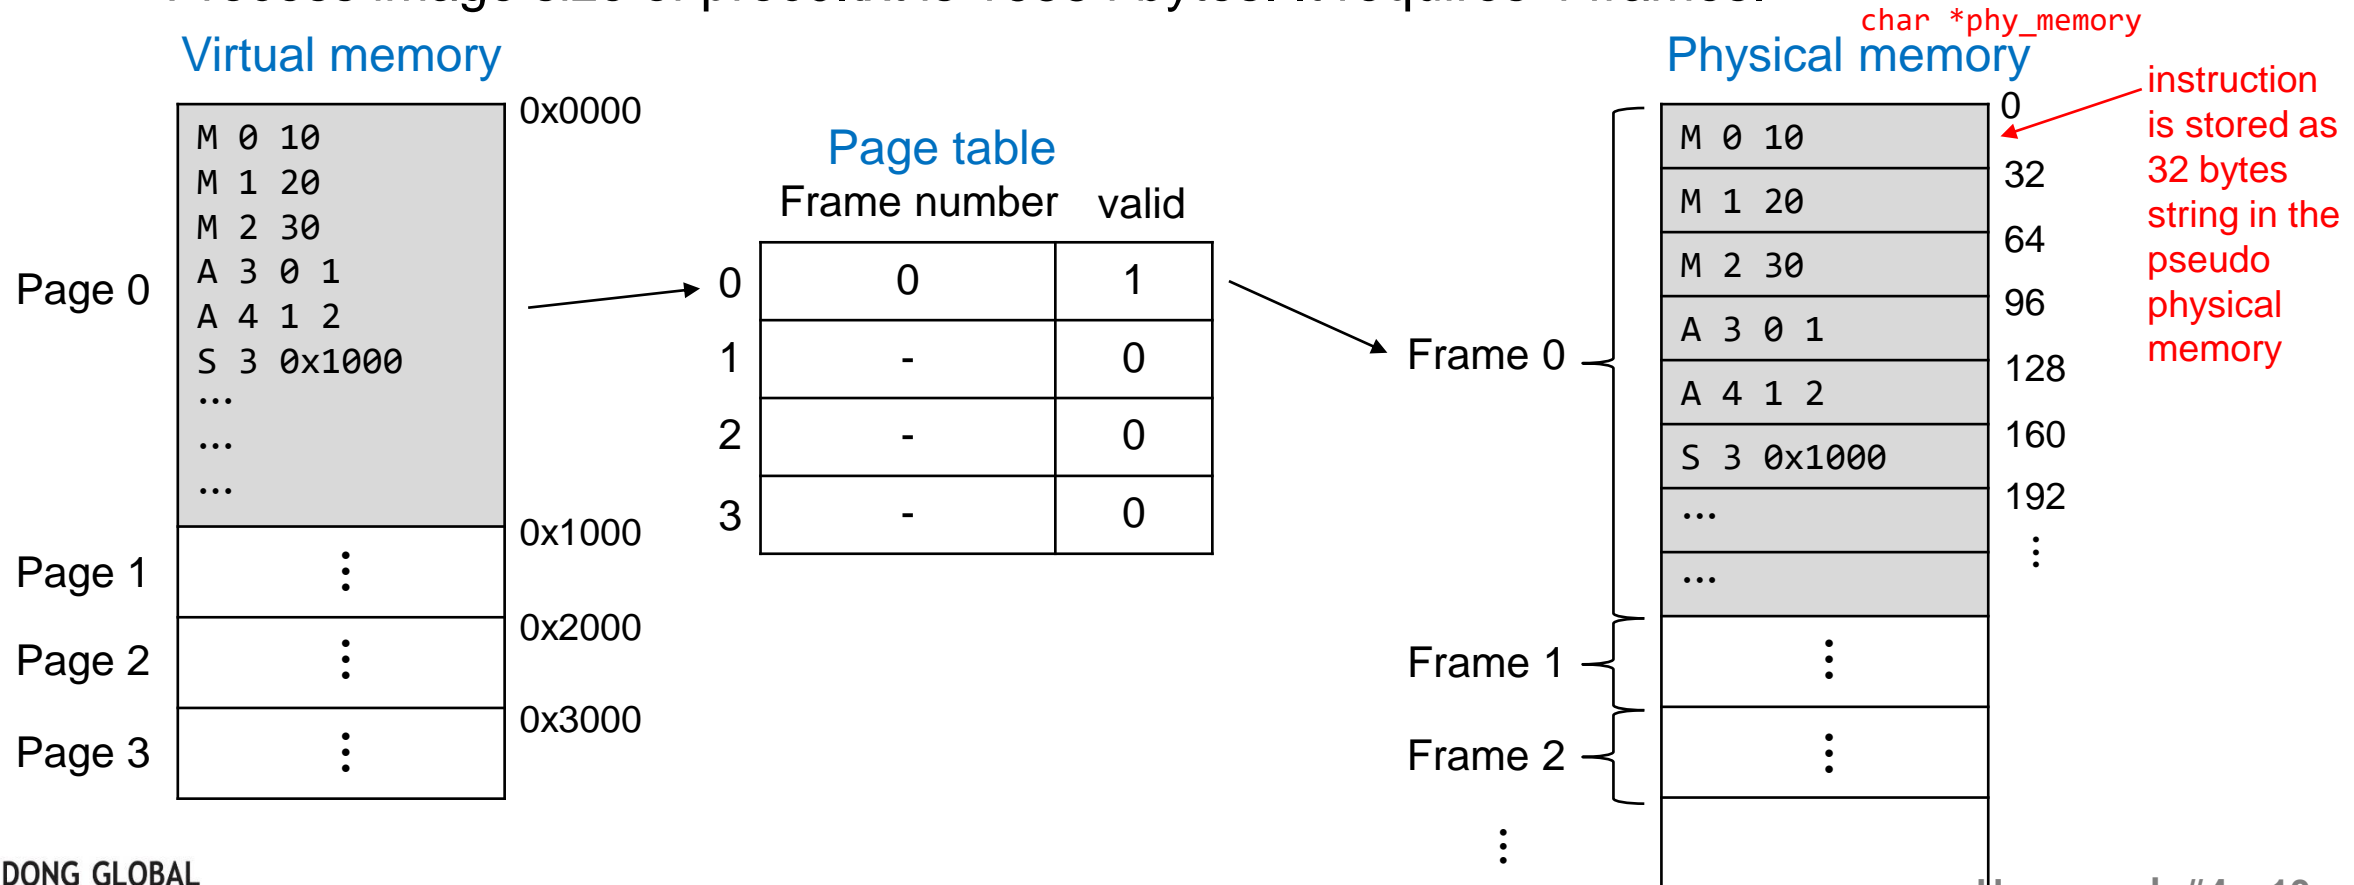
- Pseudo physical memory
 - In the simulator, physical memory is defined as 'char *phy_memory' in 'vmsim.h'.
 - The size of 'phy_memory' is defined as 'PHY_MEM_SIZE' in 'vmsim.h'.
 - The default value for PHY_MEM_SIZE is 16MB.
 - The default value for PAGE_SIZE is 4KB (4096bytes). Thus, there are 4096 frames in the 'phy_memory'.

Requirements

■ Demand paging

■ Example after 'proc0.txt' is loaded

- Process image size of proc0.txt is 16384 bytes. It requires 4 frames.



Requirements

■ Demand paging

- Example after 'proc0.txt', 'proc1.txt', and 'proc2.txt' are loaded

Virtual memory of proc0.txt

Page 0	M 0 10 M 1 20 ...
Page 1	
Page 2	
Page 3	

Page table of proc0.txt

	Frame number	valid
0	0	1
1	-	0
2	-	0
3	-	0

Virtual memory of proc1.txt

Page 0	M 0 100 M 1 10 ...
Page 1	
Page 2	
Page 3	

Page table of proc1.txt

	Frame number	valid
0	1	1
1	-	0
2	-	0
3	-	0

char *phy_memory
Physical memory

Frame 0
Frame 1
Frame 2
Frame 3
Frame 4
⋮

Page table of proc2.txt

	Frame number	valid
0	2	1
1	-	0
2	-	0
3	-	0

Virtual memory of proc2.txt

Page 0	M 4 20 M 5 60 ...
Page 1	
Page 2	
Page 3	
Page 4	
Page 5	
⋮	
Page 9	

Requirements

■ Demand paging

- Example after 'S 1 0x1000' in 'proc1.txt'

Virtual memory of proc0.txt

Page 0	M 0 10 M 1 20 ...
Page 1	
Page 2	
Page 3	

Page table of proc0.txt

	Frame number	valid
0	0	1
1	-	0
2	-	0
3	-	0

Virtual memory of proc1.txt

Page 0	M 0 100 M 1 10 ...
Page 1	110
Page 2	
Page 3	

Page table of proc1.txt

	Frame number	valid
0	1	1
1	3	0
2	-	0
3	-	0

char *phy_memory
Physical memory

Frame 0
Frame 1
Frame 2
Frame 3
Frame 4
⋮

Page table of proc2.txt

	Frame number	valid
0	2	1
1	-	0
2	-	0
3	-	0

Virtual memory of proc2.txt

Page 0	M 4 20 M 5 60 ...
Page 1	
Page 2	
Page 3	
Page 4	
Page 5	
⋮	
Page 9	

Requirements

- Demand paging
 - You don't have to implement swap in/out with backing store.
 - When page fault occurs, just allocate the page.
 - When the process has no more instructions to execute, the process terminates, and at this time, all frames allocated to the process are reclaimed.
 - The simulator finds free frame to allocate new page from pseudo physical memory using first fit scheme.
 - You don't have to implement page replacement.

Requirements

■ Page Table

- Page table entry is declared with structure 'PageTableEntry' in 'vmsim.h'.
- Each process maintains a page table defined as 'PageTableEntry *page_table' in the structure 'Process'.
- For the predefined structures please refer to 'vmsim.h'.

Requirements

■ Expected results

```
yunmin@mcn1-PowerEdge-R740:~/lecture/os/hw04/sol$ ./vmsim proc0.txt
[Clock= 0][PID=0] Opcode=M, RegNo=0, Value=10 → R[0]=10
[Clock= 1][PID=0] Opcode=M, RegNo=1, Value=20 → R[1]=20
[Clock= 2][PID=0] Opcode=M, RegNo=2, Value=30 → R[2]=30
[Clock= 3][PID=0] Opcode=A, RegDest=3, RegSrc1=0, RegSrc2=1 → R[3]=10+20=30
[Clock= 4][PID=0] Opcode=A, RegDest=4, RegSrc1=1, RegSrc2=2 → R[4]=20+30=50
[Clock= 5][PID=0] Opcode=S, RegNo=3, Addr=0x1000
[Clock= 5][PID=0] Page fault at virtual address 0x1000 (page_number=1) --> Allocated frame_number=1
[Clock= 6][PID=0] Opcode=S, RegNo=4, Addr=0x2000
[Clock= 6][PID=0] Page fault at virtual address 0x2000 (page_number=2) --> Allocated frame_number=2
[Clock= 7][PID=0] Opcode=L, RegNo=5, Addr=0x1000
[Clock= 8][PID=0] Opcode=L, RegNo=6, Addr=0x2000
[Clock= 8][PID=0] [RegisterSet]: R[0]=10, R[1]=20, R[2]=30, R[3]=30, R[4]=50, R[5]=30, R[6]=50, R[7]=0
```

※ Frame number = 0 has instructions.

→ Page fault because it is first access at 0x1000
→ R[3]=30 is stored at 0x1000 (frame number=1)
→ Page fault because it is first access at 0x2000
→ R[4]=50 is stored at 0x2000 (frame number=3)
→ R[5]=30 since 30 is loaded from 0x1000 into R[5]
→ R[6]=50 since 50 is loaded from 0x2000 into R[6]
→ Show all values of Register Set since process has no more instruction to execute

```
yunmin@mcn1-PowerEdge-R740:~/lecture/os/hw04/sol$ ./vmsim proc1.txt
[Clock= 0][PID=0] Opcode=M, RegNo=0, Value=100
[Clock= 1][PID=0] Opcode=M, RegNo=1, Value=10
[Clock= 2][PID=0] Opcode=A, RegDest=1, RegSrc1=0, RegSrc2=1
[Clock= 3][PID=0] Opcode=S, RegNo=1, Addr=0x1000
[Clock= 3][PID=0] Page fault at virtual address 0x1000 (page_number=1) --> Allocated frame_number=1
[Clock= 4][PID=0] Opcode=L, RegNo=2, Addr=0x2000
[Clock= 4][PID=0] Page fault at virtual address 0x2000 (page_number=2) --> Allocated frame_number=2
[Clock= 5][PID=0] Opcode=L, RegNo=3, Addr=0x1000
[Clock= 5][PID=0] [RegisterSet]: R[0]=100, R[1]=110, R[2]=0, R[3]=110, R[4]=0, R[5]=0, R[6]=0, R[7]=0
```

→ Show all values of Register Set since process has no more instruction to execute

Requirements

■ Expected results

```
yunmin@mcn1-PowerEdge-R740:~/lecture/os/hw04/sol$ ./vmsim proc0.txt proc1.txt proc2.txt
[Clock= 0][PID=0] Opcode=M, RegNo=0, Value=10
[Clock= 1][PID=1] Opcode=M, RegNo=0, Value=100
[Clock= 2][PID=2] Opcode=M, RegNo=4, Value=20
[Clock= 3][PID=0] Opcode=M, RegNo=1, Value=20
[Clock= 4][PID=1] Opcode=M, RegNo=1, Value=10
[Clock= 5][PID=2] Opcode=M, RegNo=5, Value=60
[Clock= 6][PID=0] Opcode=M, RegNo=2, Value=30
[Clock= 7][PID=1] Opcode=A, RegDest=1, RegSrc1=0, RegSrc2=1
[Clock= 8][PID=2] Opcode=A, RegDest=1, RegSrc1=4, RegSrc2=5
[Clock= 9][PID=0] Opcode=A, RegDest=3, RegSrc1=0, RegSrc2=1
[Clock=10][PID=1] Opcode=S, RegNo=1, Addr=0x1000
[Clock=10][PID=1] Page fault at virtual address 0x1000 (page_number=1) --> Allocated frame_number=3
[Clock=11][PID=2] Opcode=S, RegNo=1, Addr=0x3000
[Clock=11][PID=2] Page fault at virtual address 0x3000 (page_number=3) --> Allocated frame_number=4
[Clock=12][PID=0] Opcode=A, RegDest=4, RegSrc1=1, RegSrc2=2
[Clock=13][PID=1] Opcode=L, RegNo=2, Addr=0x2000
[Clock=13][PID=1] Page fault at virtual address 0x2000 (page_number=2) --> Allocated frame_number=5
[Clock=14][PID=2] Opcode=S, RegNo=5, Addr=0x2000
[Clock=14][PID=2] Page fault at virtual address 0x2000 (page_number=2) --> Allocated frame_number=6
[Clock=15][PID=0] Opcode=S, RegNo=3, Addr=0x1000
[Clock=15][PID=0] Page fault at virtual address 0x1000 (page_number=1) --> Allocated frame_number=7
[Clock=16][PID=1] Opcode=L, RegNo=3, Addr=0x1000
[Clock=16][PID=1] [RegisterSet]: R[0]=100, R[1]=110, R[2]=0, R[3]=110, R[4]=0, R[5]=0, R[6]=0, R[7]=0
[Clock=17][PID=2] Opcode=L, RegNo=1, Addr=0x3000
[Clock=18][PID=0] Opcode=S, RegNo=4, Addr=0x2000
[Clock=18][PID=0] Page fault at virtual address 0x2000 (page_number=2) --> Allocated frame_number=1
[Clock=19][PID=2] Opcode=L, RegNo=5, Addr=0x2000
[Clock=20][PID=0] Opcode=L, RegNo=5, Addr=0x1000
[Clock=21][PID=2] Opcode=S, RegNo=4, Addr=0x4000
[Clock=21][PID=2] Page fault at virtual address 0x4000 (page_number=4) --> Allocated frame_number=3
[Clock=22][PID=0] Opcode=L, RegNo=6, Addr=0x2000
[Clock=22][PID=0] [RegisterSet]: R[0]=10, R[1]=20, R[2]=30, R[3]=30, R[4]=50, R[5]=30, R[6]=50, R[7]=0
[Clock=23][PID=2] Opcode=L, RegNo=6, Addr=0x4000
[Clock=23][PID=2] [RegisterSet]: R[0]=0, R[1]=80, R[2]=0, R[3]=0, R[4]=20, R[5]=60, R[6]=20, R[7]=0
```

→ Process#1 terminated

→ Frame 1 is allocated to Process#0 using first fit scheme.
Frame 1 was released by Process#1.

→ Frame 3 is allocated to Process#2 using first fit scheme.
Frame 1 was released by Process#1.

→ Process#0 terminated

→ Process#2 terminated

Requirements

- Please check 'TODO' in 'vmsim_main.c' and 'vmsim_main.h'
- For unmentioned requirements, you can implement freely.
- Your program should be executed on Ubuntu.
- You can compile with makefile. (use 'make' or 'make clean')
- Write clean source code for readability and add proper comment in your source code
- Test your source codes with many cases for self verification
- Upload **zip** file on LMS by compressing all your source codes
 - File name: hw04_student id.zip (ex: hw04_20400022.zip)
- Due date: 11:59pm, 6/21 (Fri)